

# Resolver.jl

## A New SAT-based Resolver

*Stefan Karpinski*

<https://JuliaHub.com>

# Package Version Resolution Problem

The setup:

- There are *packages*
- Each package has *versions*
- Each version *depends* on a set of packages
- A subset *incompatible* pairs of versions
  - only makes sense for versions of different packages

# Package Version Resolution Problem

A valid solution:

- Every dependency is satisfied
- No incompatible version pairs

A solution satisfies a set of requirements (packages) when:

- The solution contains a version of each required package

# Package Version Resolution Problem

The decision problem:

- *Is there a valid solution that satisfies a set of requirements?*

This problem is somewhat famously NP-complete

- If you can solve this you can solve SAT problems
- You might as well use a SAT solver (or something stronger)

Otherwise you'll end up badly implementing a buggy, slow SAT solver

# Not Just Any Solution

Obviously just knowing if there is a solution is not that useful

- That's just how computational complexity classes are defined

# Optimal Solutions

We also don't just want any solution — we want an optimal solution

- Requires a notion of some solutions

Start by putting a preference ordering on versions of each package

- Extend this to a preference ordering on solutions
- Actually more tricky and subtle than expected

# What We Do Now

Pkg.jl includes a version resolver

- Uses belief propagation
- Heuristic: may not find solutions nor necessarily optimal ones
- But, works remarkably well

Implemented and maintained by [Carlo Baldassi](#)

- Thank you, Carlo!

# Version Numbers

The biggest issue with the existing resolver:

- Structure & meaning of version numbers deeply baked into logic
- Bakes in that higher version numbers are better
- Can't support pre-release or build numbers



# Newer $\neq$ Better?

When would you prefer older versions over newer ones?

- Version fixing — prefer current version
  - minimize changes to manifests
- Download avoidance — prefer pre-installed versions
  - avoid installing new versions if possible
- Downgrade resolution — prefer oldest allowable versions
  - useful for testing lower compat bounds

We do all of these in hacky ways currently

# Resolver.jl Approach

- Avoid coupling with details of packages, versions, registries
  - Resolver.jl doesn't know about any of these
- Use an actual SAT solver (PicoSAT)
  - how does optimization work? (you'll see)
- SAT solvers are very sensitive to problem size
  - significant preprocessing to minimize SAT problem size

**Ends up being fast & scalable**

# Emergent Features

- Since it doesn't know about version numbers
  - can give it arbitrary preference ordering
- Generates optimal solutions in lexicographical order
  - user can specify priority of packages
- Semi-internal SAT problem API is more generally useful
  - can be used for other related problems

**Also quite flexible**

# SAT

A SAT problem consists of:

- A number of variables
- A number of clauses — all must be satisfied
- Each clause is a *disjunction* of variables and negated variables

Example:  $(a \vee \neg b \vee c) \wedge (\neg a \vee b)$

- Equivalent:  $(b \Rightarrow (a \vee c)) \wedge (a \Rightarrow b)$

# Encoding Package Problems

Variables:

- One for each package:  $p$ 
  - use  $p$  and  $q$  for packages
- One for each version:  $v$ 
  - use  $v$  and  $w$  for versions

# Some Notation

Properties:

- Package of a version:  $p = P(v)$
- Versions of a package:  $v \in V(p)$
- Dependencies of a version:  $q \in D(v)$
- Set of conflicts:  $\{v, w\} \in C$
- Conflicts of a version:  $w \in C_v$

# Clauses

meaning	quantifiers	clause
version implies its package	$\forall v$	$v \Rightarrow P(v)$
package implies some version	$\forall p$	$p \Rightarrow \bigvee V(p)$
only one version of a package	$\forall p, \{v, w\} \subseteq V(p)$	$\neg v \vee \neg w$
versions imply dependencies	$\forall v, q \in D(v)$	$v \Rightarrow q$
conflicts are exclusive	$\forall \{v, w\} \in C$	$\neg v \vee \neg w$

# Code

Loop over packages

```
for p in names
    info_p = info[p]
    n_p = length(info_p.versions) # number of versions of p
    v_p = vars[p] # lookup variable for p

    # generate clauses for p
end
```



# Code

Version implies its package

```
for i = 1:n_p
    PicoSAT.add(pico, -(v_p + i)) #  $\neg(p@i)$ 
    PicoSAT.add(pico, v_p)        #  $p$ 
    PicoSAT.add(pico, 0)          # end clause
end
```

# Code

Package implies some version

```
PicoSAT.add(pico, -v_p)          #  $\neg p$ 
for i = 1:n_p
    PicoSAT.add(pico, v_p + i) #  $\neg(p@i)$ 
end
PicoSAT.add(pico, 0)             # end clause
```

# Code

Only one version of a package

```
exclusive &&  
for i = 1:n_p-1, j = i+1:n_p  
    PicoSAT.add(pico, -(v_p + i)) # ¬(p@i)  
    PicoSAT.add(pico, -(v_p + j)) # ¬(p@j)  
    PicoSAT.add(pico, 0)          # end clause  
end
```

# Code

Versions imply dependencies

```
for i = 1:n_p
    for (j, q) in enumerate(info_p.depends)
        info_p.conflicts[i, j] || continue # deps & conflicts
        PicoSAT.add(pico, -(v_p + i)) # ¬(p@i)
        PicoSAT.add(pico, vars[q])    # q
        PicoSAT.add(pico, 0)          # end clause
    end
end
```

# Code

Conflicts are exclusive

```
for i = 1:n_p
    for q in names; v_q = vars[q]
        for j = 1:length(info[q].versions)
            info_p.conflicts[i, b+j] || continue
            PicoSAT.add(pico, -(v_p + i)) # ¬(p@i)
            PicoSAT.add(pico, -(v_q + j)) # ¬(p@j)
            PicoSAT.add(pico, 0)          # end clause
        end
    end
end
```

# Preprocessing

SAT is sensitive to problem size

- Pays off to generate as small a problem as you can

Two strategies to reduce the size of problems:

1. Only include “reachable” packages and versions
2. Eliminate redundant versions that can't be picked

# Reachability

Concept:

- Start with best versions of all required packages
- Only consider next version if better version has potential conflict
- Also consider dependencies of any reachable versions

# Reachability

Recursive definition

condition	implies
$p$ in requirements	$p_1$ reachable
$p_i$ reachable and depends on $q$	$q_1$ reachable
$p_i$ reachable with reachable conflict	$p_{i+1}$ reachable
$p_i$ reachable and depends on $q$ and $q_n$ reachable with reachable conflict	$p_{i+1}$ reachable



# Redundancy

Suppose  $v$  and  $w$  are versions of the same package

- Preference:  $v > w$  ( $v$  is preferable to  $w$ )
- Dependencies:  $D(v) \subseteq D(w)$
- Conflicts:  $C_v \subseteq C_w$

There can never be a reason to choose  $w$  over  $v$

- We can delete  $w$

# Constructing a SAT instance

The process:

- Parse registry data for explicitly required packages
- Parse registry data for recursive dependencies of all versions
- Reachability analysis to find candidate package versions
- Redundancy analysis to eliminate unchoosable versions
- Construct SAT problem for remaining package versions

# Finding Optimal Solutions

SAT solvers can find some solution (or tell us there is none)

- How do we find an optimal solution?

# Finding Optimal Solutions

PicoSAT allows pushing and popping sets of clauses

- Given a solution, add clause to improve some package version
- When no longer improvable, last package version is optimal
- Pop last clause, leaves optimal solution locked in
- Repeat for each package

# Finding Multiple Solutions

There can be multiple Pareto-optimal solutions

- When an optimal solution is found
- Add clause forcing future solutions not to be dominated by it
  - In other words, must be strictly better w.r.t. some package
- Stop when no longer satisfiable
  - Or you have enough solutions

# Changing Version Preferences

This is completely straightforward now:

- The resolver doesn't care what order versions are ordered
- You just give it an ordering of versions for each package
  - Downgrade resolution — reverse normal version order
  - Download avoidance — sort pre-installed versions first
  - Version fixing — sort current version first
- This can be done on a per-package basis
  - *E.g.* Downgrade for some, normal for others

# Changing Package Priority

Solutions are optimized one package at a time

- Sorting by popularity might be a good default
- User can modify this order however they want
  - Move certain packages to the front

Can't really prioritize indirect dependencies over direct ones

- This is actually an inherent feature of the problem
- Subtle and took me a long time to understand

# Improved Error Feedback?

How to give better feedback when resolution is impossible?

- Open problem — if you have ideas, let me know!

Possible approach:

- Don't try to explain why it's impossible
- Tell people how they could fix the situation
  - Requirements they could remove
  - Conflicts they could eliminate